

# CS 111 Final exam, short

Cody S Hubbard

TOTAL POINTS

**117 / 130**

## QUESTION 1

### 1 sparse files 10 / 10

✓ - 0 pts Correct

- 2 pts Did not mention block level allocation detail for FAT

- 1 pts Did not mention pointers in Unix Systems

- 2 pts Did not demonstrate understanding of FAT table entries

- 8 pts Did not answer the question

- 3 pts Incorrect understanding of Unix System V

## QUESTION 2

### 2 Links 9 / 10

- 0 pts Correct

✓ - 1 pts Did not go into detail about name storage and datablocks

- 2 pts Did not talk about link count

- 3 pts Incorrect understanding about metadata comparison

## QUESTION 3

### 3 Distributed capabilities 7 / 10

- 0 pts Correct

✓ - 3 pts Did not go into detail about network attacks

- 3 pts Did not show understanding of capability

- 2 pts Did not explain lack of challenges in single machine environment

- 5 pts Did not answer the question

## QUESTION 4

### 4 horizontal scalability 10 / 10

✓ - 0 pts Correct

- 1 pts Missing some details.

- 3 pts Not very accurate.

- 9 pts Wrong.

## QUESTION 5

### 5 Full disk encryption 10 / 10

✓ - 0 pts Correct

- 1 pts Missing some details.

- 4 pts Not very accurate.

- 7 pts Wrong in explaining attacks.

- 9 pts Wrong.

- 10 pts Empty.

## QUESTION 6

### 6 Addressing in the OS itself 10 / 10

✓ - 0 pts Correct

- 10 pts Physical Addresses

- 5 pts The OS uses the same CPU as user processes do, and its instructions are passed through the same MMU as the instructions of those user processes.

- 5 pts Provide extra wrong answer

- 5 pts The concept is correct, but it doesn't answer the question directly.

## QUESTION 7

### 7 Disk defragmentation 10 / 10

✓ - 0 pts Correct

- 1 pts Some details missing/wrong.

- 9 pts Wrong.

- 4 pts Not accurate.

- 10 pts Empty.

## QUESTION 8

### 8 Preemptive scheduling with threads 3 / 10

- 0 pts Correct

✓ - 7 pts Voluntary yielding is not preemptive scheduling.

- 10 pts Totally wrong

- 8 pts Not very effective for real preemptive

scheduling.

- 8 pts Putting a thread into an infinite loop will definitely not preempt it.

- 8 pts The alternate interpretation is right, but then how does the library achieve preemption?

- 10 pts OS knows nothing about these threads.

- 8 pts How do you preempt a thread this way?

- 6 pts Which interrupts/system calls?

- 10 pts But how can you actual preempt?

- 10 pts This is not a problem in mutual exclusion. It concerns mechanisms for controlling execution.

- 4 pts More details on what system calls to use (i.e., timers).

- 10 pts Use timer interrupts.

- 0 pts Simpler to just ask the OS for a timer interrupt, but this probably would work, too.

- 5 pts Close. The process hosting the threads can ask for a timer interrupt and use that opportunity to switch to a different thread.

- 8 pts Without OS assistance, how could this be done?

- 0 pts Clumsy, but workable. Scheduling timer interrupts would be easier.

- 8 pts Vastly cumbersome to recast a local multithreaded program as a distributed program so you can gain control via RPC. Use timer interrupts.

- 10 pts No answer

#### QUESTION 9

### 9 Sloppy counters 10 / 10

✓ - 0 pts Correct

- 3 pts wrong answer for the first question

- 3 pts wrong answer for the second question

- 4 pts wrong answer for the third question

- 2 pts missing the key point "inaccurate" for the third question

- 10 pts wrong answer

#### QUESTION 10

### 10 Semaphores and thread initialization 8 / 10

- 0 pts Correct

- 2 pts Semaphore blocks on  $< 0$ .

- 3 pts Parent runs wait before it can run.

✓ - 2 pts Not considering both orders of operations

- 3 pts Counter should be zero

- 4 pts Post increments, wait decrements.

- 10 pts Totally wrong

- 3 pts Must specify the parent's semaphore operation.

- 6 pts No guarantee that child runs before parent, so this set of operations won't always work.

- 7 pts Not how semaphores work, as described.

- 3 pts Child just runs post, parent just runs wait.

- 3 pts Child must do initialization operations before the post().

- 7 pts As described, nobody ever gets to run. Child decrements semaphore on wait(), it's less than zero, he blocks, so nobody ever posts and wakes him (or the parent, who also blocks).

- 10 pts Semaphores, not mutexes.

- 3 pts Must specify the child's semaphore operation.

- 3 pts Must not call wait before creating child, since then there will be no one to post.

- 10 pts No answer

#### QUESTION 11

### 11 Hold-and-wait and deadlock 10 / 10

✓ - 0 pts Correct

- 2 pts While this does avoid deadlock, it doesn't guarantee progress, since the two threads can each re-acquire one of the locks, then fail to acquire the other one.

- 4 pts This is not a universal solution, since you cannot necessarily group all locks needing to be acquired together. Unless you put all locks under one lock, which serializes locking, which could be worse than deadlock. OK if you specify that the global lock can be released after acquiring other locks.

- 4 pts Very hard, in general, since many reasonable operations require more than one resource to be updated atomically, and thus locked.

- 2 pts "In advance" isn't enough. Must request them all at once.

- 10 pts Totally wrong.

- 3 pts If you seize someone's lock this way, you may run into atomicity problems.

- 3 pts If you use one universal lock to control obtaining the others, either you must hold that universal lock as long as you hold any locks, or you must specify that all locks you need are obtained at once while holding the universal lock.

- 3 pts Not allowing parallelism at all is not a great solution, and can't work for multi-process scenarios.

- 5 pts Spin locks will never resolve the deadlock.

- 2 pts Only viable if lock holders know they might lose their exclusive access. With leases, they inherently do. Otherwise, maybe not, which can cause problems.

- 5 pts Not helpful unless someone releases a lock eventually.

- 3 pts What is the condition that can't be done?

- 2 pts Not clear what you mean by "preallocating". Correct for some reasonable meanings, not correct for others.

- 4 pts Unclear explanation of why it causes deadlock.

- 4 pts Lots of reasons events don't occur. Only a deadlock if the reason is someone else holds a lock preventing it.

- 5 pts This approach attacks the mutual exclusion condition, not hold and wait. You can have mutual exclusion while still preventing hold and wait.

- 3 pts Kills parallelism if you must hold global lock while holding any other lock.

- 5 pts Priorities won't help.

- 5 pts So how do you avoid it?

- 4 pts Can happen with locks at the finest granularity.

- 5 pts Far more drastic than necessary. Without further investigation, how do you know which waiting processes to even kill?

- 2 pts How do you generalize this?

- 4 pts Not helpful to only take locks away when a

process is done with them.

- 5 pts Incorrect description of hold-and-wait problem.

- 5 pts If mutual exclusion is required on a particular piece of code, you can't move it out of lock, so this isn't a real solution.

- 5 pts Preventing preemption doesn't help.

- 5 pts Semaphores don't prevent deadlock.

- 5 pts Waking a thread waiting on a lock doesn't help.

- 5 pts Scheduling alone won't help. And not running any other process while one is holding some locks would kill performance.

- 4 pts More details on what you mean by reservation here.

- 4 pts Just checking doesn't help.

- 3 pts A bit more than that is necessary to take away a lock without causing problems.

- 5 pts Would cause serious concurrency problems.

#### QUESTION 12

### 12 Event-based concurrency and async I/O

10 / 10

✓ - 0 pts Correct

- 2 pts missing the point that there is only one thread handling all events.

- 10 pts wrong

- 2 pts missing the point that it affects performance and concurrency

- 2 pts missing the point that synchronous I/O blocks other events

#### QUESTION 13

### 13 Optimizing file writes 10 / 10

✓ - 0 pts Correct

- 3 pts Cylinder groups don't help that much on writes.

- 1 pts Delayed writes also complicates ensuring file system consistency, especially in the face of crashes.

- 2 pts Write buffering doesn't have anything to do with sequential writing. It's about delaying writes till they're cheap or unnecessary.

- **10 pts** Nothing to do with file system write

performance.

- **4 pts** Not really a write optimization. More about reads.

- **3 pts** Require more details on exactly how this optimization interacts with writes.

- **2 pts** Introduces problems with on-disk consistency.

- **10 pts** Not a write optimization

- **1 pts** Write buffering usually handled in block I/O cache.

- **3 pts** Complexities?

- **1 pts** Journaling requires garbage collection.

- **10 pts** No answer.